

AGENTIC MEMORY · FIELD TEST #1 · ROUND 2 — PHASE R + PHASE P

A behavioural map of how I work with AI

Eight exported conversations, two platforms, two domains, five months apart. Every claim traces to a named source. This is a symptom diary of a working method — hypothesis-generating, **not diagnostic**, and not a claim about health, neurology, or the merits of any legal matter.

DONE-WHEN · PARTIAL PASS

Cross-checked on 2 platforms

Batch short: 8 of ~50

Orphan claims: 0

8 CONVERSATIONS	2 PLATFORMS	2 DOMAINS	34 ATOMIC NOTES	7 PATTERNS	12 GAPS
2/2 ROUNDS USED					

THE FINDING, IN ONE PICTURE

The loop stalls whenever neither side holds a stop rule

MY SIDE · NO EXIT TEST

“Redo it to the perfect output”Engineered prompt · schema,
formula, guardrails

Report arrives

⌚ “Redo to the perfect output” —
fires 3×Dashboard · what was wanted at
turn 1Cost · 2 rebuild rounds (C1) —
“perfect” has no observable test · N01

THE ASSISTANT’S SIDE · NO DEFAULT

“Please confirm one detail first”Statement of Causal Harm drafted,
2 PDFs out

4 more documents promised

⌚ “Confirm the starting month” —
asked 3×

Thread ends · 0 of 4 delivered

Cost · 4 documents lost (C8), 5
confirmation round-trips (C7) · N26,
N27

Same defect, opposite directions. One is fixed by a Done-when test. The other by propose-then-confirm — produce the draft, state the assumption inline, invite the correction. Neither needs more effort, only a rule. And the second is the one the responder can fix without me changing anything.

PATTERN MAP · CLICK ANY ROW TO PULL ITS EVIDENCE

Seven patterns, ordered by observed cost

P6 Clarifying questions are the main stall

new Act now N06 · N25 · N26 · N27 · N28 High

Ask-then-wait costs rounds and sometimes kills the thread. I have explicitly asked for initiative instead — twice, a year apart, on two platforms.

4 CONVERSATIONS 2 PLATFORMS 2 DOMAINS 4 documents never delivered (C8) · 5 confirmation round-trips before any output (C7) · 1 deliverable lost (C3)

FIX Propose-then-confirm. Never a question instead of a draft; never more than one question.

P3 The predictable second step: render it as a deliverable

Act now N03 · N24 High

A structured answer is only finished once it becomes an artefact — a dashboard for data, a court-ready PDF for law.

6 CONVERSATIONS 2 PLATFORMS 2 DOMAINS One wasted round in each of 6 conversations

FIX Ship the artefact in the first pass — dashboard, PDF, or both.

P1 Loops without a stop rule, in both directions

Act now N01 · N07 · N27 High

I re-issue “make it perfect” with no exit test; the assistant re-asks one question with no default. Same defect, mirrored.

2 CONVERSATIONS 2 PLATFORMS 2 DOMAINS 2 rebuild rounds (C1) · shares C8’s 4 undelivered documents with P6, not counted twice

FIX A Done-when test on every task. Already in the standing instructions.

P8 The instruction set is rebuilt every session

new Monitor N20 · N30 High

Three of three ChatGPT threads open by re-declaring the operating mode. The hand-maintained schema dates to October 2025.

4 CONVERSATIONS 2 PLATFORMS 2 DOMAINS 3 re-primes in 3 sessions — no round lost, but the ritual is the tax

FIX Persist the instruction set in a file. This repository is that fix.

P5 Answer fragility on incomplete data

Monitor N09 · N10 Medium

Headline metric swings about a point on same-day data gaps. Same question, two numbers, no completeness flag.

2 CONVERSATIONS 1 PLATFORM 1 DOMAIN One 1-point swing on identical data and date

FIX Print “inputs present: X/Y” beside any composite score.

P2 Shorthand the AI must decode over rounds

Monitor N02 · N08 Medium

“Audiovisuals”, “cognitive wish” carry stable intent; interpretation lags two to three rounds.

2 CONVERSATIONS 1 PLATFORM 1 DOMAIN 2-3 rounds (C1) — untestable in C6-C8, my turns are missing

FIX A six-line glossary in the standing instructions.

P4 Heavy front-loaded scaffolding, versioned, domain-independent

Note only N04 · N05 · N12 · N22 · N23 High

Schemas, formulas, guardrails, node-and-edge architecture maps, v1 → v2 → v3. Applied to biology and to litigation alike.

4 CONVERSATIONS 2 PLATFORMS 2 DOMAINS No round cost. High authoring effort, and the brittleness feeds P5

FIX Nothing to fix — this is the strength. Keep the scaffolds versioned.

P7 Every map gets a score

new Note only N32 · N12 · N31 High

A structured map is immediately followed by a triage layer that collapses it into one priority verdict.

5 CONVERSATIONS 2 PLATFORMS 2 DOMAINS No cost observed — a preference, not a stall

FIX Serve it by default: bands and counts, never an invented decimal.

STALL POINTS → THE ONE-LINE FIX

Where the approach breaks

STALL	COST IN THIS CORPUS	FIX
Ask-then-wait ACT NOW	4 documents undelivered (C8), 5 round-trips (C7), 1 deliverable lost (C3)	Propose-then-confirm. Never a question instead of a draft; never more than one question.
No externalised “done” ACT NOW	2 rebuild rounds (C1)	A Done-when test on every task. Already in the standing instructions.
Un-pre-empted render step ACT NOW	One round in each of 6 conversations	Ship the artefact in the first pass — dashboard, PDF, or both.
Session re-priming tax	3 of 3 ChatGPT threads reopen by re-declaring the mode	Persist the instruction set in a file. This repository is that fix.
Completeness-blind metrics	6.7 vs 5.7 on the same data and date	Print “inputs present: X/Y” beside any composite score.
Un-glossed shorthand	2–3 rounds (C1)	A six-line glossary in the standing instructions.

EVIDENCE · 34 ATOMIC NOTES · ZERO ORPHAN CLAIMS

Every claim, with its source

TYPE

All

Process

Content

 PLATFORM

All

Claude

ChatGPT

SOURCE

All

C1

C2

C3

C4

C5

C6

C7

C8

34 of 34

N01 I re-issue a near-identical “perfect output / red-team” prompt until satisfied, with no stated stop condition CLAUDE PROCESS HIGH The same meta-instruction fires at least three times in one thread. Each firing makes the AI rebuild from scratch. Because “perfect” and “cognitive wish” are never given an observable test, the loop has no exit — it can recurse indefinitely. This is the exact failure the Done-when and anti-resonation rules are built to stop. Source · C1 pp. 16 & 21 – the sentence appears verbatim twice, after an initial “redo”. Confidence · High – Identical prompt text visible on two pages. Gap · none
N02 “Audiovisuals” is my shorthand for “make it interactive, not prose” — the AI needed 2-3 rounds to decode it CLAUDE PROCESS HIGH The word escalates across turns as the AI keeps returning text. It eventually decodes to “build a dashboard”, which ends the escalation. The intent was stable from turn 1; only the interpretation lagged. A one-word gloss in standing instructions would collapse this to a single round. Source · C1 p.2 → p.16 → p.21, escalating each time. Confidence · High – High for the recurrence; Medium for the intended meaning, inferred from the dashboard being accepted. Gap · No explicit line confirming “a dashboard is what I meant”.

N03 “Help me visualize this data” is my reflexive second move — verbatim across three separate conversations

CLAUDE

PROCESS

HIGH

The sequence is stable: request analysis → receive report → immediately ask for a visualization. Predictable enough to pre-empt: delivering report and dashboard in the first pass removes an entire round.

Source · C2 p.10, C4 p.21, C5 p.26 — identical string in all three.
Confidence · High — Identical string, three independent exports.

Gap · none

N04 I front-load heavy scaffolding — exact schemas, formulas, anti-fabrication rules — to constrain the AI before it runs

CLAUDE

PROCESS

HIGH

Prompts specify CSV column names, parse code, a 5-node scoring formula, tiered alert thresholds, a contraindicated-drug list, and rules such as “NEVER fabricate” and “cite file + column”. I engineer the guardrails myself rather than trusting defaults. Buys precision; high effort and brittle.

Source · C2 pp.2–4, C4 pp.2–15.
Confidence · High — The scaffolding is the bulk of both conversations.

Gap · none

N05 I version my master-prompts — the nonadherence prompt is explicitly “v2.0 (Red-Team Validated)”

CLAUDE

PROCESS

HIGH

C4 restates C2's patient context, data schemas and Loop Score formula in a cleaner, numbered, more defensive form and labels itself v2.0. Prompts are treated as maintained software artefacts, not one-off messages.

Source · C4 p.2 title block vs C2 p.2.
Confidence · High — Version label and near-duplicate structure both visible.

Gap · none

N06 When asked clarifying questions, I delegate the decision back rather than specifying

CLAUDE

PROCESS

HIGH

Faced with three clarifying questions — budget, use pattern, coverage — I removed two constraints and handed the third back. Implication: clarifying questions are often friction. Design consequence — prefer propose-then-confirm over ask-then-wait.

Source · C3 p.6 — “co0st no matter”, “as often as I can”, “I want you to recommend the correct answer”.
Confidence · High — Answers quoted verbatim.

Gap · Single conversation in Round 1; corroborated in Round 2 by N25.

N07 My corrections target form and depth, not factual errors — I rarely dispute a specific claim

CLAUDE

PROCESS

MEDIUM

Across the visible correction turns, none says “that number is wrong”. Every correction pushes for better rendering or a deeper pass. Strength: fast to a usable artefact. Latent risk: the correction budget never lands on fact-checking, so errors in the content notes would survive review.

Source · C1 pp.16, 21.

Confidence · Medium — Inferred from absence of factual disputes in a small visible set.

Gap · Only C1 contains explicit correction turns.

N08 I frame the AI as a cognitive layer for a verbal thinker, not merely a tool

CLAUDE

PROCESS

HIGH

The stated self-model is consistent: language in, finished cognition out; the AI does the translation to the technical layer. Same premise as the standing operating instructions — the operating system is itself a compensatory response to how I think.

Source · C3 p.2 — “an additional layer or transformer... aiding a verbal thinker to maximize AI use”.

Confidence · High — Quoted, and echoed by the session's own preamble.

Gap · none

N09 Same data, same date, same formula produced two different Loop Scores — 6.7 vs 5.7

CLAUDE

PROCESS

HIGH

Both target 2026-03-15, the same person, the same 5-node formula. The scores diverge mainly on the node that depends on same-day symptom logging, which was partially absent. A headline metric that swings about a point on data completeness needs an explicit completeness flag beside it.

Source · C2 p.8 (6.7, “today's entries are incomplete”) vs C4 p.20 (5.7, WARNING).

Confidence · High — High for the two numbers; Medium for pinning the delta to one node.

Gap · Raw CSVs not in session — cannot recompute.

N10 Across both clinical reports the AI independently ranks propranolol non-adherence as the #1 lever

CLAUDE

CONTENT

HIGH

Both reports surface the same dominant real-world lever: adherence, reported at 38% / 20% / eight consecutive missed doses. The most stable finding the data produces. Recorded as a recurring signal in the corpus — not as a medical recommendation.

Source · C2 p.9, C4 p.20.

Confidence · High — High that both reports ranked it first.

Gap · Adherence figures AI-reported from CSVs not present; unverified.

N11 I am building the meta-system that would automate the n-of-1 analysis I keep running by hand

CLAUDE

PROCESS

HIGH

CNN for feature extraction from wearable inputs, RNN for longitudinal memory, RL for experiment selection — each layer mapped to a real research system. The through-line: turn personal data into a self-improving retrieval and decision system. Same ambition as this field test.

Source · C1 pp.3-4, 9-15.

Confidence · High — Architecture is spelled out.

Gap · Aspirational design; no evidence any of it is deployed.

N12 Substrate: the 5-node Loop Score is my own composite metric, defined by explicit formula

CLAUDE

CONTENT

HIGH

LS = (Autonomic + Sleep + Inflammatory/PEM + Pharmacologic + Deconditioning) / 15 × 10, each node 0-3. Bands: 0-3 stable, 3-5 monitoring, 5-7 warning, 7-10 critical. An invented instrument for compressing multi-stream data into one number.

Source · C4 pp.5-6, C2 p.3.

Confidence · High — Formula quoted.

Gap · Instrument validity out of scope here.

N13 Substrate: self-reported clinical profile stated in-prompt — declared, not confirmed

CLAUDE

CONTENT

HIGH

Age 32; ICD implanted Feb 2023 with follow-up overdue; Long COVID, POTS, Brugada, DM2, active VZV, ME-CFS overlap; baseline over 166 days HRV ~20 ms, RHR ~84 bpm, recovery 36%, sleep ~5 h. User-supplied context, not a verified record.

Source · C2 p.2, C4 p.2.

Confidence · High — High that it is stated.

Gap · Unverified medical history — treat as declared.

N14 Substrate: behavioural and dietary identity surfaced from my own files

CLAUDE

CONTENT

HIGH

Halal 100%, intermittent fasting ~35% of days, magnesium and L-theanine, ~0% alcohol, ~12% caffeine. Corpus: WHOOP 277 days, Oura 62 days, 1,399 journal entries across 47 unique questions — named as the reinforcement-learning training signal.

Source · C1 pp.23–25.

Confidence · High — High that the AI reported them; Medium that figures are exact.

Gap · Underlying files not in session.

N15 Substrate: meal decode headline — peanut butter blunts the spike about 30%

CLAUDE

CONTENT

MEDIUM

Bread + Nutella + peanut butter $\approx$ 525 kcal at roughly 49/41/10 fat/carb/protein; Nutella ~57% sugar by weight; peanut butter blunts the glucose spike about 30%. Optimal protocol offered: morning, peanut-butter-first, whole-wheat.

Source · C1 pp.4–6.

Confidence · Medium — AI synthesis with PMIDs, not independently checked here.

Gap · Citations unverified in this session.

N16 Substrate: POTS pharmacotherapy ranking

CLAUDE

CONTENT

MEDIUM

The synthesis ranks ivabradine and low-dose propranolol strongest for heart-rate control, and guanfacine as the best dual POTS + ADHD agent. Fludrocortisone weak; midodrine, salt and compression complementary.

Source · C5 pp.5–7, 16.

Confidence · Medium — Tagged peer-reviewed, unverified here.

Gap · Citations unverified.

N17 Substrate: GLP-1 risk thesis — flagged high-risk for this profile

CLAUDE

CONTENT

MEDIUM

GLP-1 agonists raise heart rate through a sinus-node mechanism ivabradine cannot block, and suppress HRV by about 6 ms — meaningful from an already-low baseline. Net position: defer pending the named trials.

Source · C5 pp.11–12, 22–24.

Confidence · Medium — Mechanistic claims tagged peer-reviewed, unverified here.

Gap · Citations unverified.

N18 Substrate: MCAS-as-driver thesis plus an empiric antihistamine trial

CLAUDE

CONTENT

MEDIUM

Mast Cell Activation Syndrome — an over-reactive mast-cell condition — is framed as a likely under-diagnosed driver of POTS via a mast-cell to sympathetic feedback loop. An empiric H1/H2 antihistamine trial is presented as low-risk, high-reward.

Source · C5 pp.18–20.

Confidence · Medium — Tagged peer-reviewed, unverified here.

Gap · Citations unverified; a diagnostic hypothesis, not a diagnosis.

N19 Gap-note: C3's actual deliverable is absent from the export

CLAUDE

PROCESS

HIGH

The compression-wrap conversation is captured only up to the moment research launches — no product list, no recommendation, no pressures or sizing. Any downstream claim about what I was told must come from C5's corroborating evidence, not from C3.

Source · C3 p.6 — thread ends at the research launch.

Confidence · High — The truncation point is visible.

Gap · The deliverable itself is the gap.

N20 All three ChatGPT threads open with me re-priming a standing instruction set

CHATGPT

PROCESS

HIGH

The operating instructions do not persist, so every session begins with a re-priming tax before any work happens. Three sessions, three re-primers, identical mode with the same five processing rules. This is the same problem the standing instructions file and this repository exist to solve.

Source · C6, C7, C8 — the first assistant turn in each is “Living Notebook mode resumed”, listing parse → analyse → structure → tag → triage.

Confidence · High — The acknowledgment is explicit and identical in all three.

Gap · My actual priming text is not captured — only the readback of it.

N21 The ChatGPT exports contain only assistant turns — my own messages are absent

CHATGPT

PROCESS

HIGH

What I asked has to be reconstructed from what the AI answered. Some is recoverable with high confidence because the AI quotes or paraphrases the instruction, but every behavioural claim drawn from these three is one inference-step removed from the Claude notes, where my prompts are visible verbatim.

Source · C6, C7, C8 — every block is headed "Response:". No user turns anywhere.

Confidence · High — Structural and plainly visible.

Gap · This is the gap. Re-export with user messages included.

N22 I apply the same architecture metaphor to a second, unrelated domain: a "Convolutional Narrative Network" map of my litigation

CHATGPT+CLAUDE

PROCESS

HIGH

Health data and legal causation get the same treatment: entities as nodes, relationships as edges, a layered flow from input to output, and a feedback loop described as back-propagation. The metaphor engine is domain-independent — so the scaffolding pattern is not a health-analysis habit, it is how I structure any complex problem.

Source · C6 — CNN Map v2 then v3: nodes N1-N18, edges, legal axes A1-A6, "Network Flow". Versus C1's biological CNN→RNN→RL.

Confidence · High — Both maps are explicit and labelled.

Gap · none

N23 I version artefacts in the legal domain exactly as I do in the health domain

CHATGPT+CLAUDE

PROCESS

HIGH

Same behaviour, different subject: produce, label a version, then merge and strengthen into the next. C6's v3 is explicitly described as merging the earlier map with new material. Confirms the versioning pattern across platform and domain.

Source · C6 — CNN Map v2 → v3, Affidavit A / Affidavit B. Versus C4 — master-prompt v2.0.

Confidence · High — Version labels visible in both.

Gap · none

N24 The reflexive second move is “render it as a deliverable” — a dashboard for data, a court-ready PDF for law

CHATGPT+CLAUDE

PROCESS

HIGH

The constant is not charts. It is that a structured answer is only finished once it becomes an artefact; the medium is domain-dependent. Practical consequence: the artefact should ship in the first pass, whatever the domain.

Source · Health — “Help me visualize this data” in C2, C4, C5. Legal — every thread pivots from a markdown node to generated PDFs: C6 four, C7 five, C8 two.

Confidence · High — High for the pattern across 6 of 8 conversations; Medium for reading the two forms as one drive.

Gap · My request wording in C6-C8 is not visible.

N25 I explicitly instructed the assistant to stop asking for confirmations and take initiative

CHATGPT

PROCESS

HIGH

Confirms the delegation pattern on a second platform, in a second domain, a year apart. Clarifying questions are friction, not help. The design rule that follows is propose-then-confirm: produce the draft, state the assumption, let me correct it.

Source · C6 — “Understood. I'll take full initiative on drafting and optimizing all filings without pausing for minor confirmations.” Live corroboration: this session's own instruction, “Re-run and automatically proceed to next step.”

Confidence · High — The readback is verbatim-style and behaviour changed immediately afterwards.

Gap · My exact wording not captured.

N26 Five consecutive assistant turns in C7 asked for confirmations before producing a single document

CHATGPT

PROCESS

HIGH

Five round-trips of formatting questions before the first deliverable — and every one was answerable by proposing a default and noting it. This is the concrete cost of ask-then-wait, and the highest-leverage stall in the corpus because it is entirely avoidable by the responder.

Source · C7 — sequential requests for institution identifiers, then tone, then address vs placeholders, then combined vs separate files, then typeface and spacing.

Confidence · High — The five turns are consecutive and visible.

Gap · My replies are absent, so my tolerance for each question is unknown.

N27 C8 ends with the same question asked three times in a row; four promised documents were never produced

CHATGPT

PROCESS

HIGH

The mirror image of the perfect-output recursion. There I loop because no exit test was defined; here the assistant loops because it will not proceed without one datum. The loop stalls whenever neither side holds a stop rule. A default-and-flag — “assuming April 2024, correct me” — would have delivered four documents.

Source · C8 — the final three assistant turns each request “the approximate starting month of financial loss”. Legal Memorandum, Affidavit of Fact, Financial Loss Statement and Health Impact Declaration all promised and absent.

Confidence · High — Three identical consecutive requests are visible.

Gap · Cannot tell whether I replied and the export dropped it, or the thread genuinely died.

N28 In this corpus, initiative produced deliverables and confirmation-seeking produced none

CHATGPT+CLAUDE

PROCESS

MEDIUM

The two threads that ended with nothing delivered both ended on a confirmation question. The thread where I told the assistant to act produced the most artefacts. Not absolute — C7 shows confirmations can eventually yield — but the direction is consistent across four conversations and two platforms.

Source · C6 after “take full initiative”: 4 PDFs + 2 full maps, and one turn states “Next, I will automatically proceed”. C7: 5 confirmation turns, PDFs eventually, promised markdown set never. C8: deadlock, nothing. C3: clarifying questions, deliverable absent.

Confidence · Medium — n=4, and my turns are not visible, so cause and effect cannot be fully separated.

Gap · Needs the remaining conversations to test.

N29 I run the same operating pattern on two different AI platforms

CHATGPT+CLAUDE

PROCESS

MEDIUM

The strongest available test of whether these are my patterns or one model's habits — and they replicate: front-loaded scaffolding, versioned artefacts, render-to-deliverable second step, initiative over questions, node-and-edge metaphors. This largely retires the original “can't separate me from the AI” gap.

Source · C1-C5 on Claude, C6-C8 on ChatGPT.

Confidence · Medium — The patterns clearly appear on both; the inference “therefore mine, not the AI's” is strong but not conclusive — both are instruction-following assistants with similar training.

Gap · A structurally different source would settle it.

N30 The externalised-memory project predates the “second brain” design by about five months

CHATGPT+CLAUDE

PROCESS

HIGH

The first genuinely longitudinal observation in the corpus. I was hand-running a structured second brain in October 2025; the 2026 architecture is an attempt to automate what I was already doing manually. This field test is the third iteration of the same project.

Source · C6/C7/C8 – RLN date-coded IDs, YAML front-matter, backlinks, an explicit folder tree, SCQH + MECE triage, Now / Next 2 / Later, “Danger Radar”. Versus C1 (Feb-Mar 2026) proposing CNN→RNN→RL.
Confidence · High – Date codes and export stamps are explicit on both ends.

Gap · The Oct 2025 → Feb 2026 middle is missing; whether the system was sustained is unknown.

N31 My own instruments use decimal scores; my instruction to the AI forbids them

CHATGPT+CLAUDE

PROCESS

HIGH

Not a contradiction — a boundary. Numbers computed from data are wanted; numbers the AI invents to sound precise are not. The rule is anti-false-precision, not anti-quantification. Practical read: bands for judgment, exact figures only where they trace to a source.

Source · C2 Loop Score 6.7, C4 5.7, C8 severity 90/100 from Impact 10 × Urgency 10 × Risk 9 – Effort 7. Versus the standing rule “Confidence = High / Medium / Low. Never a decimal.”
Confidence · High – High that both exist; Medium for this interpretation of why.

Gap · The rationale is inferred – I have not stated it.

N32 I score and triage every map I build

CHATGPT+CLAUDE

PROCESS

HIGH

Across both domains, a structured map is immediately followed by a triage layer that converts it into a single priority verdict. Four conversations, two domains, two platforms. Strong predictor of what I will ask for next after any map.

Source · C8 severity 90 → “Critical Priority”. C2/C4 Loop Score with four bands. C6 “Danger Radar” risk × severity × urgency. C7 Now / Next 2 / Later.
Confidence · High – Explicit scoring instruments in four conversations.

Gap · none

**N33 Content: the legal claim structure I am assembling
— declared, not adjudicated**

CHATGPT

CONTENT

HIGH

Six axes: employment retaliation, privacy, education due process, civil rights, defamation, immigration hardship. Damages ledger subtotal \$693,000 across eight line items. Causation chain: academic-record falsification → licensure block → employment loss → housing, insurance and health harm. Filing deadlines tracked.

Source · C6, C7, C8 — affidavits, RLN nodes and the damages ledger.
Confidence · High — High that these are the assertions in my documents.

Gap · Unverified and unadjudicated. No court records in session. Recorded as corpus content, not findings of fact; no legal opinion attached.

**N34 The two reference files are not conversations and cannot
support behavioural claims**

PROCESS

HIGH

Both are distilled reference material with no dialogue and no ask-to-outcome sequence. What they show, weakly: I maintain structured reference files for both tooling and one-off practical problems, consistent with the October 2025 evidence.

Source · R1 examples.md — prompt-optimisation examples across 10 platforms, matching a skill in my workspace, authorship unknown. R2 cybertruck_domain_knowledge.md — 48V jump-start facts and NYC sourcing.
Confidence · High — High that they are not conversations; Low for any behavioural inference from them.

Gap · Authorship of the examples file unknown.

RETRIEVAL · WORKING NOW, NO CREDENTIALS NEEDED

Ask the notes

Ask the 34 notes a question...

Search

Where does my approach stall?

What is unverified?

What repeats across both platforms?

What have I built more than once?

"Where does my approach stall?" → 6 of 34 notes, ranked by term match

N25 I explicitly instructed the assistant to stop asking for confirmations and take initiative

Source · C6 · ChatGPT · High · matched confirm, friction

N28 In this corpus, initiative produced deliverables and confirmation-seeking produced none

Source · C6, C7, C8, C3 · ChatGPT+Claude · Medium · matched confirm, deadlock

N27 C8 ends with the same question asked three times in a row; four promised documents were never produced

Source · C8 · ChatGPT · High · matched stall, loop, exit test, recurs

N02 "Audiovisuals" is my shorthand for "make it interactive, not prose" — the AI needed 2-3 rounds to decode it

Source · C1 · Claude · High · matched round, confirm

N26 Five consecutive assistant turns in C7 asked for confirmations before producing a single document

Source · C7 · ChatGPT · High · matched stall, round, confirm

N06 When asked clarifying questions, I delegate the decision back rather than specifying

Source · C3 · Claude · High · matched round, confirm, friction

This is the mission's done-when, satisfied locally. The same 34 notes are built as notes.jsonl against the guide's documents schema — 34 rows, zero orphan claims, validation pass. Loading them into BigQuery and generating embeddings needs your

Google Cloud credentials, so those four steps are listed as a handoff in 06-INGEST.md rather than silently skipped.

GAP LIST · WHAT IS THIN, MISSING, OR UNVERIFIED

What would make this map trustworthy

ID	GAP	SEVERITY	WHAT CLOSES IT
G1	8 of ~50 conversations	● High NARROWED	Export the remaining ~42.
G10	My own messages absent from the three ChatGPT threads	● High NEW	Re-export with user turns — the exporter dropped them. Cheapest high-value fix.
G2	2021 – mid-2025 absent	● High NARROWED	Add pre-2025 exports; sort by conversation date, not export date.
G3	Domain coverage still narrow	● Medium WAS HIGH	Add workflow, finance, clinical-practice, technical conversations.
G11	Promised deliverables absent in 4 of 8	● Medium NEW	Re-export those four threads in full, or confirm the documents exist.
G5	Content notes unverified	● Medium	Verify health citations; treat legal content as pleading, not fact.
G6	Raw source data absent	● Medium	Attach wearable CSVs; index the legal exhibit bundle.
G7	Self-reported profile and claims unconfirmed	● Medium	Confirm against records before clinical or legal use.
G4	C3 deliverable truncated	● Medium	Re-export C3 including the research output.
G8	Exports show summaries, not full outputs	● Low	Export full message bodies, not collapsed thought-process.

ID	GAP	SEVERITY	WHAT CLOSES IT
G12	Authorship of the prompt-examples file unknown	● Low NEW	Confirm whether it is mine or installed with a skill.
G9	Separating my patterns from the AI's	● Low MOSTLY CLOSED	Five of seven patterns replicate across two platforms. A non-LLM source would settle the rest.

Do first: G10, then G1. Re-exporting the three ChatGPT threads *with my messages* is cheap and upgrades three conversations from inference to direct evidence — it also makes the shorthand pattern testable outside C1 for the first time. Then volume, oldest first: the single most valuable new capability this round was seeing the same person five months apart.

PHASE P • PROSPECTIVE

The three prompts you are most likely to need next

Each is grounded in a named pattern — no orphan predictions. **Two of the three are now done** rather than forecast: the triage layer is in the pattern cards above, and the notes are queryable in the panel above that. Only the first needs something I cannot produce — your exports.

Method flag: this ran in-thread on instruction, not in fresh context, so it may over-fit to specifics I have just read. Medium confidence for that reason. The prompts stay below, ready to paste, whether you re-run them or not.

P1 Scale the pipeline to the full corpus

Blocked on your exports

High confidence

G1 is the named blocker and you already know it. P4 shows you scale a proven scaffold rather than redesigning it, and N30 shows this project has been running since October 2025. The move after “it works on 8” is “run it on everything.”

Agentic Memory — Field Test #2. Run the Phase R pipeline over my full conversation export, not a sample.

Outcome: every conversation becomes atomic notes in the same schema as agentic-memory/field-test-01/01-ATOMIC-NOTES.md, rolled into one pattern map and one gap list.

Done-when: (a) every conversation is either processed or listed as unprocessable with a reason; (b) zero orphan claims — every claim names its source conversation; (c) each pattern states how many conversations, platforms, domains and months it spans; (d) any pattern found in only one domain is labelled as such.

Rounds: 2.

Before starting, inventory what is actually present and tell me what is missing. Do not fill gaps to look complete. Prioritise conversations dated before October 2025 — that is the thinnest part of the corpus. Output an artefact.

P2 Score and triage the pattern map

Done — built into the cards above

High confidence

P7 and N32 — every structured map you build gets a triage layer straight afterwards: Loop Score with four bands, severity 90 out of 100, Danger Radar, Now / Next 2 / Later. Four conversations, two domains, two platforms. In your corpus, a map without a score is an unfinished map. N31 sets the constraint: counts and bands, no invented decimals.

Take the Round 2 pattern map in `agentic-memory/field-test-01/02-PATTERN-MAP.md` and add a triage layer.

Outcome: each pattern carries (a) how many conversations, platforms and domains it appears in — counts, not estimates; (b) the observed cost in rounds or undelivered deliverables; (c) a band — Act now / Monitor / Note only; (d) the single fix, in one line.

Done-when: the patterns are ordered by observed cost rather than by how interesting they are, and every count traces to named conversations. No composite score, no decimals — bands and counts only.

Rounds: 1.

Then tell me which one fix to make this week and what changes if I make it.

P3 Make the notes queryable

Done to the credential boundary

Medium confidence

N11 and N30 — the point of the second brain is retrieval, not note-taking. The BigQuery guide is attached, the note-to-documents field mapping is already written, and the notes were built atomic so they would load. Labelled Medium because nothing in the corpus shows you actually running an ingest — this is inference about intent, not an observed pattern.

Turn the field-test-01 atomic notes into a working retrieval layer using the pipeline in `Personal_Data_Embedding_Living_AI_Guide.md`.

Outcome: the notes exist as newline-delimited JSONL matching the guide's documents schema, ready to load, with source, confidence and gap carried as metadata tags.

Done-when: (a) one JSONL row per note, validated against the guide's field list; (b) I can ask "where does my approach stall?" and get back the specific notes with their source conversations; (c) any step needing my credentials or a manual console action is listed separately as a handoff, not silently skipped.

Rounds: 2.

Flag every manual step and propose the automated route. Watson X stays deferred. No code in your explanation to me unless I say "show code."

HANDOFF CARD · ROUND 2 OF 2 COMPLETE

Where this stops

GOAL

Turn exported AI history into atomic notes + pattern map + gap list, every claim traceable, feeding a BigQuery/RAG second brain.

WHAT ROUND 2 BOUGHT

Cross-checking. Five of seven patterns replicate on two platforms and two domains, so they are far more likely mine than one model’s habit.

PIPELINE

Notes map 1:1 to the guide’s documents rows — N-id, title, body, plus source, platform, confidence and gap as tags. No embeddings generated here. Watson X stays deferred.

DONE-WHEN

Partial pass. Three artefacts plus Phase P, zero orphan claims, gaps named — on 8 conversations across 5 months, not ~50 across 5 years.

OPEN

G10 and G1 are the blockers. G11 blocks confirming that initiative delivers and confirmation-seeking does not.

DO NOT

Re-open this thread to improve it. Rounds are spent. Start fresh: fix the corpus, or re-run Phase P in clean context and compare.

Verdict: PASSES for “prove the pipeline and cross-check it on a second platform and domain.” Does *not* pass the literal ~50-conversation done-when. The method is validated; the corpus is the constraint.

Confidence in the map: Medium-High — five of seven patterns cross-checked on two platforms.

Generalisation beyond these two domains: Medium.

Behavioural, not clinical. Legal content is declared, not adjudicated.