

SESSION ANATOMY

One Claude Code session, measured: 13 things he typed, 261 turns, 121 tool calls, 13 files produced. Built so a third party can analyse it and re-create it — and so the parts that *cannot* be recovered are stated rather than quietly filled in.

7 characters →

15,769 tokens of reasoning · 25 tool calls

Exchange 10 was the word **“Do next”**. It produced nine reasoning blocks, twenty-five tool calls, two published files and a working ingest pipeline. **He saw about 4% of what was generated.** That asymmetry — small input, large hidden work, small visible output — is the shape of every productive exchange in this session.

13 HIS MESSAGES	261 TURNS	344 BLOCKS	121 TOOL CALLS	60 REASONING BLOCKS	~94k HIDDEN TOKENS
13 FILES MADE	PASS ROUND-TRIP				

THE LIMIT, STATED FIRST

MY REASONING TEXT IS NOT IN THIS EXPORT

Not redacted here — never written. All 60 reasoning blocks were stored **signature-only**.

What survives for each reasoning block: that it happened, where in the conversation, its signature hash and length, and roughly how many tokens it consumed.

What does not: a single character of what it said. No amount of re-processing the transcript changes this — the plaintext was never on disk.


Anyone reconstructing “the model’s thinking” from this data is reconstructing its **shape and cost**, not its content. A claim otherwise is false, and this export gives you the means to check it: `SELECT * FROM blocks WHERE kind='thinking'` returns 60 rows with `chars = 0` and `redacted = 1`.

HIS THINKING VS MINE · ONE ROW PER MESSAGE HE TYPED

THE LEDGER

“Hidden” is reasoning he never saw. “Shown” is the reply he read. Both bars are scaled to the largest exchange, so the row heights are comparable across the session.

#	WHAT HE TYPED	CHARS IN	REASONING	HIDDEN TOK	HIDDEN	REPLY CHARS	TOO
1	/model – switch to claude-opus-5 · <i>pre- compaction, work not recorded</i>	143	–	–		–	
2	(command output) Set model to claude-opus-5 · <i>pre- compaction, work not recorded</i>	71	–	–		–	
3	Session continuation + prior- context summary · <i>pre- compaction, work not recorded</i>	22,780	–	–		–	
4	/compact – with standing output preferences · <i>pre- compaction, work not recorded</i>	1,767	–	–		–	
5	(compacted) – resume work	55	8	14,703		129	
6	[interrupted] continue	42	1	189		–	

#	WHAT HE TYPED	CHARS IN	REASONING	HIDDEN TOK	HIDDEN	REPLY CHARS	TOT
7	[interrupted] continue	42	1	165		—	
8	[interrupted] continue	42	1	654		1,476	
9	5 new files attached – “Re-run and automatically proceed”	481	12	32,463		3,153	
10	“Do next”	7	9	15,769		2,642	
11	“What’s next?”	13	2	1,468		1,945	
12	“Go and write it.”	16	6	4,905		2,875	
13	/spine /mega- loop – export everything	424	19	22,741		1,084	

Exchanges 1-4 show input with no work attached. That is not a parsing failure — the session was compacted, and the assistant records from before the compaction boundary did not survive. The work happened; the evidence of it did not. Stated rather than hidden, because a reader counting rows would otherwise conclude the session began with four ignored messages.

WHERE THE EFFORT WENT

INSTRUMENT PROFILE

TOOL CALLS · 121 TOTAL, 7 ERRORS	
Bash	43
Edit	42
Write	12
Read	11
Artifact	8
SendUserFile	3
Skill	2
Errors recovered from	7

PRIVACY ZONES · EVERY BLOCK TAGGED	
Confidential — litigation	11
Confidential — medical	13
Private work — identifier	10
Private work — internal	310
Public	0
Zone ceiling	Confidential

FIDELITY CHECK	
turns	261 = 261
blocks	344 = 344
tool_calls	121 = 121
usage	117 = 117
exchanges	13 = 13
artifacts	13 = 13
SQL → live DB → JSON	PASS

SESSION FACTS	
Models	opus-4-8, opus-5
His input	25,883 chars
His replies read	13,304 chars
Reasoning (est)	93,861 tok
Tool output returned	1.4M chars
Files produced	13, checksummed

 DERIVATIONS · CHECK THESE BEFORE TRUSTING ANY NUMBER

TWO FIGURES ARE ESTIMATES

```
hidden_tokens_est = output_tokens - (visible_chars ÷ 4)
```

output_tokens	exact, billed. Covers reasoning + reply + tool arguments.
visible_chars	exact. Everything actually emitted.
÷ 4	the standard chars-per-token approximation. NOT a measurement. The only soft term in the calculation.

```
shown_fraction = visible_tokens ÷ (visible + hidden). Inherits the same approximation.
```

Everything else is exact — character counts, block counts, tool counts, timestamps, SHA-256 checksums. If you have a real tokeniser, substitute it for the ÷ 4 and both estimates improve. The builder is included so you can.

 STRUCTURE

SEVEN TABLES, FOUR VIEWS

turns → blocks is the spine; everything else hangs off it.

```

session — turns — blocks — tool_calls (joined on tool_use_id)
          |           |
          |           └─ usage
          └─ exchanges one message he typed + everything it caused
          └─ artifacts  files produced, with SHA-256

block      one content unit: text | thinking | tool_use | tool_result | plain
turn       one logical message. Records sharing a message id are grouped —
           without that, usage counts once per block instead of per response
           and every derived figure inflates.
exchange   the unit of "his thinking vs mine"
```

```

-- The finding this export exists to show
SELECT exchange_no, ask_chars, hidden_tokens_est, reply_chars, shown_fraction
FROM    v_his_vs_mine
ORDER BY hidden_tokens_per_ask_token DESC;

-- Prove the reasoning is unrecoverable rather than taking my word
SELECT COUNT(*), SUM(chars), SUM(redacted) FROM blocks WHERE kind='thinking';
```

THE FOLDER

WHAT IS WHERE

FILE	SIZE	PURPOSE
session.json	2.2 MB	The full export. Every table, every row, every surviving character. See here.
session.db	2.1 MB	SQL already loaded. Query immediately, no setup.
session.sql	2.1 MB	Portable DDL + INSERTs, SQLite dialect.
session.md	274 KB	Readable transcript in order, reasoning blocks marked unrecoverable
README.md	—	Schema, derivations, limits, re-creation steps.
MANIFEST.json	3 KB	Counts, totals, per-file SHA-256, stated limits.
VERIFY.txt	—	The mechanical grade.
build_export.py	31 KB	Regenerates all of it from the raw transcript.

To re-create the session: read CLAUDE.md for the rules it ran under, replay the ask column of exchanges in order, then compare output against the 13 checksummed artifacts. Expect the reasoning to differ — you have the inputs and the outputs, not the thinking between them. A faithful re-creation reproduces the **artifacts**, not the thought process.

Before sharing this onward: the Confidential blocks hold his own legal filings and clinical history. This page carries structure and measurements only — no verbatim Confidential payload. The full-fidelity files do carry it, and stay local.

Derived from the raw transcript by build_export.py. Nothing summarised, inferred, or filled in.
Round-trip verified: PASS.
Interpretation unverified — by contract, the builder cannot grade itself.